

# COMPUTER VISION

## 2018 - 2019

### >IMAGE FEATURES

UTRECHT UNIVERSITY

RONALD POPPE

# IMAGE CLASSIFICATION

**Aim is to find regions in an image that correspond to objects of interest**

**We will first look at how to represent (regions in) an image**

- Lectures 5-6

**We then address the training and testing of machine learning classifiers**

- Lectures 7-8

**Finally, we turn to convolutional neural nets for detection**

- Lectures 9-11

# OUTLINE

**Image descriptors**

**Applications**

**Issues**

**Low-level image descriptors**

**Histograms of oriented gradients (HOG)**

**Scale-invariant feature transforms (SIFT)**

# IMAGE DESCRIPTORS

# IMAGE DESCRIPTORS

**Describe the characteristics of an image:**

- Derived from the pixels

**Describe an image (or part of it) in a compact way**

- Should ideally be invariant to nuisance factors (viewpoint, scale, illumination, etc.)

**Similar images should have similar image descriptors**

- As usual: what is similar?

# IMAGE DESCRIPTORS<sup>2</sup>

**Different object, different image**





# IMAGE DESCRIPTORS<sup>3</sup>

**Same object, different image**



# IMAGE DESCRIPTORS<sup>4</sup>

**Same object, same image?**





# IMAGE DESCRIPTORS<sup>5</sup>

## **Simplest image descriptor: the pixels in an image**

- Width \* height \* channels number of dimensions
- Not really compact
- Not really invariant to nuisance factors



# IMAGE DESCRIPTORS<sup>6</sup>

**Image descriptors are often termed image features**

- A feature vector is a vector representation with each number/dimension derived from the image
- E.g.  $\mathbf{x} = (x_1 \dots x_n)$  with  $n = \text{rows} * \text{columns} * \text{color channels}$  and each dimension is a pixel color value

**Today, we will look at various image descriptors:**

- What are they used for?
- Which nuisance factors should they be invariant to?
- How to calculate them?

# APPLICATIONS

# APPLICATIONS

**Image stitching**

**Object detection**

**Duplicate detection**

**Video stabilization**

# IMAGE STITCHING

**When images partly overlap, they can be stitched together**


- Useful for making panoramas
- Used in Google Street View, any smartphone and camera

**Overlap is never perfect. Differences in:**

- Rotation
- Scale
- Lighting
- Etc.

# IMAGE STITCHING<sup>2</sup>

**Example: stitch two images together**





# OBJECT DETECTION

## Object detection is the process of finding objects in an image

## Two types of object detection

- Generic: recognize classes of objects (all cars)
- Specific: recognize a specific instance (BMW 118)



# OBJECT DETECTION<sup>2</sup>

## Different levels of granularity:

- Just saying if the image depicts the object: recognition
- Finding the location (bounding box): detection
- Determining which pixels are part of it: segmentation





# DUPLICATE DETECTION

**Duplicate detection considers the problem of finding images that are identical apart from:**

- Scale (resolution)
- Framing (crop region)
- Encoding (jpeg)
- Coloring (grayscale/rgb, variations, copies)
- Rotation (only in 2D)

**This is a verification task:**

- Are these two images near-duplicates?

# DUPLICATE DETECTION<sup>2</sup>

## Typical examples:

- Find pictures of the Mona Lisa
- Find the same picture but in a larger resolution
- Find out who uses your (copyrighted) pictures



# FACE VERIFICATION

**Face verification is the process of determining whether a face belongs to a specific user**

- This is (obviously) a verification task

**Can be done based on**

- 2D images
- 2D + depth images (Kinect)
- Full 3D model (range scanner)
- Near infra-red

# FACE VERIFICATION<sup>2</sup>

**Applications are mainly in security**


- Login with your face instead of password
- To filter out people with a neighborhood restriction



# FACE VERIFICATION<sup>3</sup>

**But applications can also be outside security:**

- Searching for pictures of (famous) people
- Understanding/recognizing movies/series
- Automatic subtitle generation



# ISSUES

# ISSUES

**The same object or scene can appear differently in images**

- Viewpoint, illumination and image quality affect the image

**We don't want all factors to influence the image descriptors**

- We call such factors nuisance factors
- It's often favorable to have an image descriptor that is invariant to many nuisance factors



# VIEWPOINT

**Images can be taken from the same viewpoint, but with a different rotation**

- This is called in-plane rotation






# AFFINE PROJECTION

Or from the same distance but with a different angle

- This is termed an affine transformation



# PERSPECTIVE PROJECTION

Or from another distance, with objects further away being smaller

- This is termed a perspective transformation



# DIRECT LIGHTING

**The direction of the light causes variation in**

- Shadows
- Specular highlights



# INDIRECT LIGHTING

Indirect lighting refers to (the amount of) ambient light

Less light lowers the contrast

- Color values are in smaller space



# IMAGE QUALITY

**The image quality is determined by:**

- Image compression
- Resolution
- Color depth

# IMAGE COMPRESSION

**Images typically stored with compression**

- JPEG, GIF, PNG, etc.

**Compression can take many forms**

- Usually detail is lost
- “noise” or “patterns” can be introduced




Original Image



GIF without dithering




GIF with dithering



# IMAGE RESOLUTION

**When reducing the resolution of an image**

- Details get lost
- Contrast regions become less pronounced



# IMAGE RESOLUTION<sup>2</sup>

What is depicted here?





# OBJECT ARTICULATION

**Objects are not always rigid but often consist of parts that can move**

- We call these objects articulated
- Detecting them is more difficult as they have different “shapes”
- Typical for humans in action




# RECAP

**When describing and comparing images, we would like our representations to be invariant to these issues**

- How to describe images so we can achieve this?

**We distinguish between low-level and high-level image descriptors**

- Low-level descriptors are close to the pixels
- High-level descriptors are more semantic, on a higher abstraction level



# QUESTIONS SO FAR?

# LOW-LEVEL IMAGE DESCRIPTORS

# LOW-LEVEL IMAGE DESCRIPTORS

**Low-level image descriptors operate on pixels of an image**

- Also called local descriptors

**The most common ones are based on:**

- Color (intensity)
- Edges (contrast)
- Motion (only for video, discussed in next lecture)

# LOW-LEVEL IMAGE DESCRIPTORS<sup>2</sup>

**Makes sense for object detection:**

- Objects typically stand out from their surroundings by different colors
- These also cause high contrast values



# COLOR DESCRIPTORS

**We looked at those before:**

- Mean color
- Color histogram (equidistant bins, from clusters)
- Gaussian mixture model


**Several “flavors”:**

- Different color spaces (RGB, HSV, etc.)
- Per channel, or combined

# COLOR DESCRIPTORS<sup>2</sup>

## Color descriptors for object recognition

- Can be a good cue
- Can be a bad cue





# EDGES

**Edges arise when neighboring pixels have contrasting intensities**

- Each pixel can be an edge pixel or not



# EDGES<sup>2</sup>

**Calculated by taking the derivative of a pixel in both the horizontal and vertical direction**

- Edges have a direction (orientation) and magnitude (strength)

**Invariant to:**

- Specific color

# PIXEL DERIVATIVES

**We can take the derivative of a pixel by applying filters:**

- $G_x$  is derivative in horizontal direction
- $G_y$  is derivative in vertical direction


**Prewitt:**  $G_x = \begin{bmatrix} -1 & 0 & +1 \\ -1 & 0 & +1 \\ -1 & 0 & +1 \end{bmatrix} * A$  and  $G_y = \begin{bmatrix} +1 & +1 & +1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix} * A$

**Sobel:**  $G_x = \begin{bmatrix} +1 & 0 & -1 \\ +2 & 0 & -2 \\ +1 & 0 & -1 \end{bmatrix} * A$  and  $G_y = \begin{bmatrix} +1 & +2 & +1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix} * A$

# PIXEL DERIVATIVES<sup>2</sup>

We apply filters using convolution:

- Center pixel is replaced by weighted sum of filter and image
- Calculated using dot product



For patch **A** and 3x3 kernel **G**:

- $A'(x, y) = \sum_{i=-1 \dots 1, j=-1 \dots 1} A(x + i, y + j) * G(x + i, y + j)$

# PIXEL DERIVATIVES<sup>3</sup>

**Example:**

-1	0	1
-1	0	1
-1	0	1

 $*$ 

1	1	0	0	0
1	1	1	0	0
1	1	1	1	0
1	1	1	1	1
1	1	1	1	0

 $=$ 

	-1	-2	-2	
	0	-1	-2	
	0	0	-2	

The diagram illustrates a 3x3 pixel derivative calculation. The input is a 3x3 grid of values:  $\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$ . This is multiplied (indicated by  $*$ ) by a 5x5 kernel. The kernel is a 5x5 grid of values:  $\begin{bmatrix} 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 1 & 1 & 1 & 1 & 0 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & 1 & 0 \end{bmatrix}$ . The result is a 5x5 grid of values:  $\begin{bmatrix} & & & & \\ & -1 & -2 & -2 & \\ & 0 & -1 & -2 & \\ & 0 & 0 & -2 & \\ & & & & \end{bmatrix}$ . Red boxes highlight the 3x3 region of the kernel used for the calculation (rows 2-4, columns 1-3) and the corresponding 3x3 region of the result (rows 2-4, columns 2-4).

# PIXEL DERIVATIVES<sup>4</sup>

Once we have the derivatives in x- and y-direction, we calculate the gradient magnitude as follows:  $G = \sqrt{G_x^2 + G_y^2}$


And the gradient orientation:  $\Theta = \text{atan2}(G_y, G_x)$




# PIXEL DERIVATIVES<sup>5</sup>

**Gradient magnitude and direction are informative:**

- Magnitude is indicator of contrast
- Direction determines the direction of the edge







**Direction can be noisy when magnitude is low**



# PIXEL DERIVATIVES<sup>6</sup>

To get a binary edge image, we can put a threshold on the gradient magnitude

- Noisy pixels typically have strong edges
- Determining threshold is subjective







# PIXEL DERIVATIVES<sup>7</sup>

## Different ways of suppressing noisy pixels:

- First apply a Gaussian filter (convolution)

## Each pixel becomes weighted average of surrounding pixels

- Image appears more blurry
- Variance can be tuned (sigma)
- For discrete values: different window sizes (5 x 5)



$1/256 \times$

1	4	6	4	1
4	16	24	16	4
6	24	36	24	6
4	16	24	16	4
1	4	6	4	1

# PIXEL DERIVATIVES<sup>8</sup>

**When applying a Gaussian filter:**

- Noisy pixels are averaged
- Details are also lost (is this a problem?)

**Typically, edge detection is preceded by Gaussian filtering with a small window size (3 x 3, 5 x 5)**



Original Image



Blurred Image ( $\sigma=10$ )



Blurred Image ( $\sigma=20$ )



Blurred Image ( $\sigma=40$ )

# CANNY EDGE DETECTION

**Another way to suppress noise due to noisy pixels:**

- Look at neighboring pixels

**Canny edge detection uses the following steps:**

- Gaussian filtering (small window)
- Obtain gradient magnitude/direction per pixel (Sobel, Prewitt)
- Non-maximum suppression
- Tracing edges

# CANNY EDGE DETECTION<sup>2</sup>

**We want edges at points with (locally) maximum magnitude**

- We want to ignore edges with lower values
- Non-maximum suppression is an edge thinning technique (remember erosion)

**Consider the two neighbors orthogonal to the gradient direction**

- Suppress the current pixel if it hasn't the maximum magnitude

**Four options:**

- Left – right
- Top – bottom
- Left-top – right-bottom
- Right-top – left-bottom



# CANNY EDGE DETECTION<sup>3</sup>

**Eventually, we are only interested in edges that are “longer”**


- Short edges considered to correspond to noise

**Canny edge detection uses two-step approach:**

- First filter out edges with low gradient magnitude (threshold)
- Next, trace each edge in the direction orthogonal to the gradient (allow some flexibility in direction)
- Only pixels belonging to longer traces are kept

# CANNY EDGE DETECTION<sup>4</sup>

**Result of Canny edge detection based on Prewitt operator**



# RECAP

## **Color and edges say something about an image**

- Objects stand out from their surroundings based on color or contrast

## **Neither color descriptors nor edges invariant to:**

- Rotation
- Scale
- Viewpoint
- Location in the image

**We need better image descriptors!**

# QUESTIONS?



# HISTOGRAM OF ORIENTED GRADIENTS

# HOG

**Histograms of oriented gradients (HOG) densely encode edges within a grid [Dalal & Triggs, CVPR 2005]**

**The representation is invariant to scale**


- Based on bounding box

**HOGs are somewhat invariant to**


- Rotation (small angles)
- Local variations
- Illumination and color (based on edges)

# HOG<sup>2</sup>

Can you guess which object was on the picture of which this HOG descriptor was calculated:



And this one?

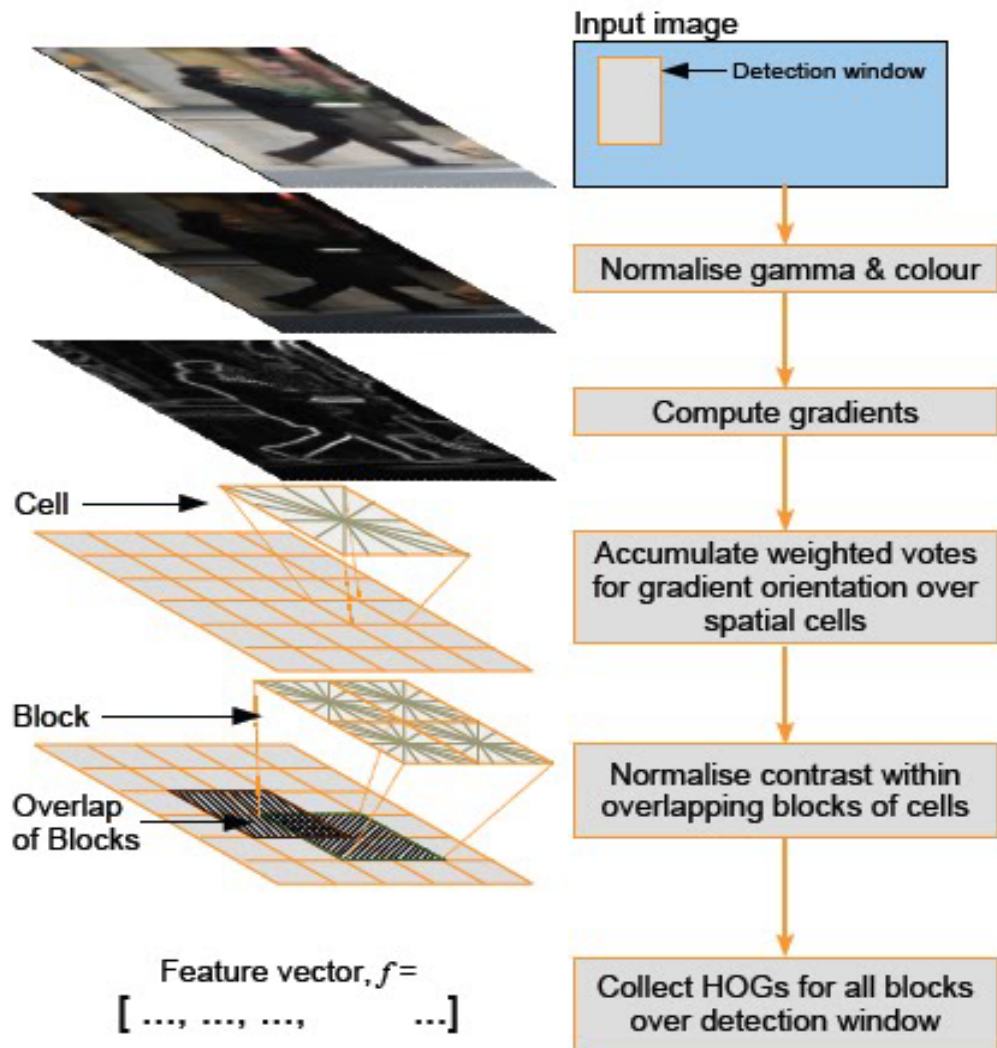


# HOG<sup>3</sup>

**HOGs are calculated in several steps, based on a bounding box:**

- Apply Gaussian filter + normalize color
  - Calculate edge magnitude and orientation
- } Edge detection!
- Summarize edges for all pixels per cell in the grid
  - Summarize blocks of cells
  - Normalize descriptor to unit length

# HOG<sup>4</sup>



# HOG<sup>5</sup>

**Grids contain  $n \times m$  cells, usually of equal size**

- So number of pixels per cell is also equal

**Blocks contain several cells, typically  $2 \times 2$**

- Blocks are overlapping
- Each block is normalized to account for local intensity differences



# HOG<sup>6</sup>

**Per cell, we construct a histogram of orientations**

- Bins can be 0-20, 20-40 degrees, etc., or 0-45, 45-90, etc.

**Each pixel's orientation contributes to orientation histogram of the cell**

- Bins determined by gradient orientation
- The amount of “weight” determined by gradient magnitude




# HOG<sup>7</sup>

**Next step is to calculate blocks:**

- Histograms in a block of neighboring cells are concatenated
- Length of the concatenated histogram is normalized to unit length (sum = 1)

**Example: block of 4 x 4 cells, each 4 x 4 pixels, and 8 orientations**





# HOG<sup>8</sup>

## **Advantages of HOG:**

- Can be calculated quickly (edge derivatives calculated once)
- Quite robust to local variations (especially within the cell)
- Quite robust to illumination changes (due to block normalization)

## **HOGs were introduced for pedestrian detection:**

- Given a patch (region in the image), determine if it is a person
- Patch described as a HOG descriptor
- Each HOG descriptor was then classified as corresponding to a person or not

# SCALE-INVARIANT FEATURE TRANSFORM

# SIFT

**Scale-invariant feature transform (SIFT) is an algorithm to describe image features [Lowe, 1999]**

- It is commonly used in matching, stitching and tracking
- Try out the demo: <http://www.cs.ubc.ca/~lowe/keypoints/>

**SIFT features are invariant to:**

- Scale
- Rotation
- Partially to viewpoint changes
- Partially to illumination changes

# SIFT<sup>2</sup>

## **Similarities with HOG:**

- Basis are gradient differences
- Final descriptor shares similarities with orientation histogram

## **Differences with HOG:**

- Not calculated at each pixel but only at specific points (sparse vs. dense)
- Encodes scale and rotation
- Can deal with partial occlusion

# SIFT<sup>3</sup>

**The SIFT algorithm has several steps:**

- Detect scale-space extrema
- Detect keypoints
- Determine orientation
- Determine local descriptor

**Once the local descriptors have been determined, they can be used for matching**

# SIFT<sup>4</sup>

**First step is to find locations and scales that can be repeatedly assigned under different views of the same object**

- Remember the nuisance factors!

**When the distance to an object changes, so does:**

- The size in the image
- The amount of detail that is visible

**SIFT addresses these issues using:**

- Size: pyramid images
- Detail: Gaussian filters

# SIFT<sup>5</sup>


**SIFT takes an image and analyses it at different scales**

- Each scale is half the previous one
- All images together form a pyramid

**At each level, Gaussian filters are applied**

- Different levels of variance

**Cope with objects of different sizes and sharpness**



# SIFT<sup>6</sup>

**Example with 3 scales and 6 levels of Gaussian filtering**






# SIFT<sup>7</sup>

**Images with the same scale but different Gaussian filtering are compared pairwise:**


- This is termed a Difference of Gaussian (DOG)

**Larger differences correspond to pixels that differ from their surroundings**

- These locations are interesting
- Typically edges and corners



# SIFT<sup>8</sup>



# SIFT<sup>9</sup>


Once we have the DOG for different scales, we need to select the local minima/maxima:

- Coarser scales are interpolated

**Compare each pixel to:**

- Its 8 neighbors on the same level
- Its 9 neighbors from scale above
- Its 9 neighbors from scale below

**Pixel is selected if it is the maximum**



# SIFT<sup>10</sup>

**We get many keypoints, some of which are unstable**


**We can refine the selection:**

- Remove keypoints with low contrast
- Remove keypoints along an edge

**We use gradient orientation and magnitude again**



**Remaining keypoints are at corners**

- Repeatable!



# SIFT<sup>11</sup>

**After selection:**



# SIFT<sup>12</sup>

## We now have a set of keypoints

- Each has a location
- Each has an orientation (determined from gradient)
- Each has a scale (determined from scale in which maximum was found)




# SIFT<sup>13</sup>

**We can calculate the local descriptor:**

- Histogram of orientations in a grid around the keypoint, with scale and orientation taken into account

**Very similar to HOG. Small differences:**

- Orientations are weighted with a Gaussian centered on the keypoint (pixels further away have less influence)
- Normalization is slightly different



# SIFT<sup>14</sup>


**SIFT points can be matched based on Euclidian or Chi squared distance**

**For object matching, we want to find pairs of matching SIFT features:**

- First find candidate features
- Then look at the distance, orientation and scales between the pairs
- Do some filtering and set a threshold on the matches
- The more matching pairs, the better the object match



# SIFT<sup>15</sup>



# RECAP

**We looked in-depth at two state-of-the-art techniques**

**Histograms of oriented gradients are fast and can be matched densely in an image**

- Ideal for pedestrian detection, but also for objects with similar orientation

**Scale-invariant feature transforms are somewhat slower but can cope with differences in viewpoint, illumination and scale**

- Ideal of object detection if the orientation of the object can vary

# QUESTIONS?

# ASSIGNMENT

# ASSIGNMENT

## **Assignment 3:**

- Tracking based on color models
- Don't start too late, this assignment is more elaborate!

**Deadline is Sunday March 10, 23:00**

**Assignment help session Thursday 11:00-12:45, RUPPERT-042**

# NEXT LECTURE

## **Optical flow**

- Motion in video
- Also an image feature

**Tuesday 13:15-15:00, BESTUURS-LIEREGG**

# QUESTIONS?